

AI-DRIVEN HIRING AND RESUME SCREENING

CASE STUDY REPORT

submitted by

ANTONE JOSEPH (VML24AD034)

MINHA NAJEEB P V (VML24AD077)

NIYATHI NITHIYANANDAN (VML24AD087)

PARVATHI KUNIYIL KATTLE (VML24AD089)

PAULBIN CHACKO (VML24AD090)

As part of the Case Study under Continuous Internal Evaluation in
the course

PEADT412 – Data Science Privacy & Ethics



Vimal Jyothi Engineering College, Chemperi
(Januvary 2026)

DECLARATION

We undersigned hereby declare that the case study report entitled "**AI-Driven Hiring and Resume Screening**" submitted as part of the Case Study under Continuous Internal Evaluation for the course **PEADT412 – Data Science Privacy & Ethics** is a bonafide work carried out by us.

This submission represents our original work and ideas expressed in our own words. Wherever ideas or words of others have been included, they have been properly cited and referenced. We further declare that we have adhered to the ethics of academic honesty and integrity and that this report has not been submitted previously, in part or in full, for the award of any degree, diploma or title at any University or Institution.

We understand that any violation of the above declaration may result in disciplinary action as per the rules of the institution and the University.

Place: **CHEMPERI**

Date: **10/03/2025**

Name & Signature of Members

**VIMAL JYOTHI ENGINEERING COLLEGE,
CHEMPERI**

CERTIFICATE

This is to certify that the case study report entitled "**AI-Driven Hiring and Resume Screening**" submitted by **Antone Joseph (VML24AD034),Minha Najeeb P V (VML24AD077),Niyathi Nithiyanandan (VML24AD087),Parvathi Kuniyil Kattle (VML24AD089),Paulbin Chacko (VML24AD090)** in partial fulfillment of the requirements for the Case Study under Continuous Internal Evaluation for the course **PEADT412 – Data Science Privacy&Ethics** is a bonafide record of work carried out by them during the academic year 2026. This report has not been submitted to any other University or Institute for the award of any degree or diploma.

FACULTY-IN-CHARGE

HEAD OF DEPARTMENT

ABSTRACT

The rapid advancement of Artificial Intelligence (AI) has significantly transformed traditional recruitment and hiring processes, particularly in the area of resume screening and candidate shortlisting. With the increasing adoption of online job portals and digital recruitment platforms, organizations often receive a large volume of applications for a single job role. Manual screening of resumes under such conditions is time-consuming, inconsistent, and highly susceptible to human bias, fatigue, and subjective judgment. As a result, many organizations are turning towards AI-driven hiring systems to improve efficiency, accuracy, and scalability in recruitment.

AI-driven hiring and resume screening systems leverage machine learning, natural language processing, and data analytics techniques to automatically extract information from resumes, evaluate candidate qualifications, and rank applicants based on job relevance. These systems aim to reduce time-to-hire, enhance consistency in decision-making, and assist recruiters by providing data-driven insights. However, despite their potential benefits, such systems raise serious concerns related to data privacy, algorithmic bias, transparency, accountability, and ethical decision-making.

This case study presents a detailed examination of an AI-based resume screening system used in employment contexts. It explores the overall architecture of the system, including data collection, preprocessing, feature extraction, model training, evaluation, and deployment. Special emphasis is given to ethical and privacy considerations, such as handling sensitive personal data, ensuring fairness across demographic groups, preventing discriminatory outcomes, and complying with data protection regulations. The study also discusses the role of explainability and human oversight in maintaining trust and accountability in automated hiring systems.

Furthermore, this case study analyzes the advantages and limitations of AI-driven recruitment tools and highlights potential risks associated with their misuse. Based on the analysis, the report proposes best practices and recommendations for the responsible and ethical deployment of AI in hiring. The findings of this study aim to contribute to a better understanding of how AI can be effectively integrated into recruitment processes while respecting privacy, promoting fairness, and upholding ethical standards in employment decision-making.

Contents

1	Introduction	7
2	Background and Motivation	8
3	Problem Statement	8
4	Objectives of the Study	9
5	Literature Review	9
6	Existing Hiring Systems	10
7	AI Techniques Used in Resume Screening	11
7.1	Natural Language Processing (NLP)	11
7.2	Machine Learning Models	11
7.3	Deep Learning Approaches	11
7.4	Fairness and Bias Mitigation Techniques	11
8	System Architecture of AI-Driven Hiring	11
8.1	Architecture Overview	12
9	Ethical Issues in AI-Driven Hiring	12
10	Privacy Concerns in Resume Screening	13
11	Bias and Discrimination in Automated Hiring	14
12	Legal and Regulatory Considerations	14
13	Risk Analysis and Challenges	15
14	Mitigation Strategies and Best Practices	15
15	Results and Discussion	15
16	Advantages of AI-Driven Hiring	16
17	Limitations of AI-Driven Hiring	17
18	Comparative Analysis of Hiring Approaches	17
19	Conclusion	18

20 Future Scope	19
21 References	20

LIST OF FIGURES

- Fig. 1 High-level architecture of an AI-driven hiring and resume screening system
- Fig. 2 AI recruitment outcomes and Results
- Fig. 3 Future trends in AI-driven recruitment systems

LIST OF TABLES

- Table 1 Comparison between Traditional Hiring and AI-Driven Hiring

1 Introduction

Recruitment is one of the most critical functions of any organization, as it directly influences workforce quality, organizational culture, productivity, and long-term success. The process of hiring suitable candidates involves evaluating not only technical qualifications but also experience, adaptability, communication skills, and cultural compatibility. Traditionally, recruitment has relied on manual resume screening, face-to-face interviews, and human judgment. While these methods have been effective for small-scale hiring, they become inefficient and inconsistent when applied to large applicant pools. With the rapid growth of digital platforms, online job portals, and professional networking websites, organizations today receive an overwhelming number of job applications for a single position. Recruiters often have limited time to review each resume, sometimes spending only a few seconds per application. This increases the risk of overlooking qualified candidates and introduces inconsistencies due to human fatigue, unconscious bias, and subjective decision-making. As a result, traditional hiring methods struggle to scale effectively in modern recruitment environments.

Artificial Intelligence (AI) has emerged as a transformative technology capable of addressing these challenges. AI-driven hiring and resume screening systems leverage machine learning algorithms, natural language processing (NLP), and data analytics to automate the early stages of recruitment. These systems are designed to parse resumes, extract relevant information such as skills, education, and work experience, and rank candidates based on their suitability for a given job role. By automating repetitive tasks, AI systems aim to reduce time-to-hire, improve efficiency, and assist recruiters in making data-driven decisions. AI-based hiring systems also offer the advantage of consistency. Unlike human evaluators, AI models apply the same evaluation criteria uniformly across all candidates, potentially reducing random variation in screening outcomes. Advanced models are capable of semantic analysis, allowing them to understand the contextual meaning of resume content rather than relying solely on keyword matching. This enables more accurate matching between job descriptions and candidate profiles, especially for technical and interdisciplinary roles.

This case study focuses on an AI-driven resume screening system implemented in an employment context. It examines the technical architecture, data processing methods, and decision-making mechanisms of the system, while placing strong emphasis on privacy, ethics, fairness, and accountability. By analyzing both the benefits and risks associated with AI-driven hiring, this study aims to provide a comprehensive understanding of how such systems can be responsibly designed and deployed in real-world recruitment scenarios.

2 Background and Motivation

The motivation for adopting AI in hiring originates from the increasing complexity of recruitment processes. Modern organizations often receive thousands of resumes for a single job opening, especially in technical and graduate-level roles. Manual screening under such conditions is time-consuming, expensive, and inconsistent.

One of the primary motivations is efficiency. AI systems can process large volumes of resumes within seconds, extracting relevant information such as education, skills, work experience, and certifications. This allows recruiters to focus their time on interviews and strategic decision-making rather than administrative tasks.

Another motivation is consistency. Human screening decisions can vary significantly due to fatigue, cognitive bias, or subjective judgment. AI systems, when properly trained, apply the same criteria uniformly across all applicants, reducing random variation in screening outcomes.

Additionally, AI-driven systems provide analytical insights that traditional hiring methods cannot easily deliver. These include identifying skill gaps in the applicant pool, tracking hiring funnel metrics, and predicting candidate success probabilities. Such insights help organizations refine job descriptions and recruitment strategies.

However, motivation alone is insufficient justification for adoption. AI systems must be designed responsibly, with mechanisms to detect bias, explain decisions, and ensure compliance with employment laws and data protection regulations.

3 Problem Statement

Organizations face several challenges in traditional resume screening and hiring processes. These challenges form the core problem addressed by this case study.

Firstly, manual resume screening does not scale well with high application volumes. Recruiters may spend only a few seconds per resume, increasing the risk of overlooking qualified candidates. Secondly, subjective human judgment introduces inconsistency and unconscious bias into decision-making. Thirdly, organizations lack systematic methods to evaluate and audit hiring decisions over time.

The problem is further complicated by legal and ethical constraints. Employers must ensure non-discriminatory hiring practices, protect applicant data, and provide transparency where required by law. Failure to address these issues can result in reputational damage and legal consequences.

This case study addresses the problem of designing an AI-driven resume screening system that balances efficiency and fairness while maintaining transparency, accountability, and legal compliance.

4 Objectives of the Study

The main objectives of this case study are as follows:

- To understand the role of artificial intelligence in automating resume screening and candidate shortlisting.
- To analyze the architecture and workflow of an AI-driven hiring system.
- To evaluate the advantages and limitations of AI-based recruitment tools.
- To examine ethical, legal, and social issues associated with algorithmic hiring.
- To propose best practices for responsible deployment of AI in employment contexts.

These objectives guide the structure and analysis presented in subsequent sections of the report.

5 Literature Review

The application of Artificial Intelligence in recruitment and hiring has gained significant attention in both academic research and industry practice. Researchers have explored how machine learning algorithms can assist recruiters by automating resume screening, predicting candidate performance, and reducing hiring costs. Early studies focused on rule-based expert systems that filtered resumes using predefined keywords and thresholds. While effective to some extent, these systems lacked flexibility and often failed to capture the semantic meaning of candidate profiles.

Recent literature emphasizes the use of Natural Language Processing (NLP) techniques to understand unstructured resume text. NLP enables systems to extract relevant entities such as skills, education, work experience, and certifications from resumes written in diverse formats. Transformer-based models, such as BERT, have further improved semantic understanding by capturing contextual relationships between words, allowing better matching between job descriptions and resumes.

However, several studies highlight ethical challenges associated with AI-driven hiring. Research has shown that models trained on historical hiring data may inherit biases related to gender, ethnicity, age, or educational background. These biases can result in

discriminatory outcomes if not properly mitigated. Scholars recommend fairness-aware learning, bias audits, and human-in-the-loop systems as essential safeguards.

From a privacy perspective, literature stresses the importance of handling sensitive personal data responsibly. Resumes often contain personally identifiable information (PII), making compliance with data protection regulations crucial. Transparent data usage policies and consent mechanisms are commonly recommended.

Overall, existing research supports the use of AI in recruitment while strongly emphasizing the need for ethical governance, transparency, and continuous monitoring.

6 Existing Hiring Systems

Traditional hiring systems rely heavily on manual resume screening performed by recruiters or HR professionals. In such systems, resumes are reviewed individually, often under strict time constraints. This approach is not scalable for large organizations and may lead to inconsistent decision-making due to human fatigue and cognitive bias.

Applicant Tracking Systems (ATS) represent a more structured approach to recruitment. ATS platforms store candidate data, track applications, and allow recruiters to filter resumes using keywords and basic criteria such as education level or years of experience. While ATS improves organization and workflow management, it still relies largely on keyword matching, which can overlook qualified candidates whose resumes use unconventional terminology.

Modern AI-powered hiring platforms extend ATS functionality by integrating machine learning models. These systems automatically parse resumes, extract structured features, and rank candidates based on predicted relevance. Some platforms also incorporate online assessments, video interview analysis, and psychometric evaluations to provide a more holistic view of candidates.

Despite technological advancements, existing systems face challenges related to fairness, transparency, and explainability. Many commercial tools operate as “black boxes,” offering limited insight into how decisions are made. This lack of transparency raises ethical and legal concerns, particularly in regulated employment environments.

7 AI Techniques Used in Resume Screening

AI-driven resume screening systems employ a combination of techniques from machine learning, natural language processing, and data analytics.

7.1 Natural Language Processing (NLP)

NLP is used to process unstructured resume text. Techniques such as tokenization, named entity recognition, and part-of-speech tagging help identify key information including skills, job titles, organizations, and durations of employment. Semantic similarity measures allow matching resumes with job descriptions beyond simple keyword overlap.

7.2 Machine Learning Models

Supervised learning models are commonly used to predict candidate suitability. These models are trained on historical hiring data labeled with outcomes such as shortlisted, interviewed, or hired. Algorithms like logistic regression, decision trees, and gradient-boosted models are popular due to their interpretability and efficiency.

7.3 Deep Learning Approaches

Deep learning models, especially transformer-based architectures, are increasingly used for resume-job matching. These models learn dense vector representations of text and can capture contextual meaning, improving matching accuracy for diverse resumes. However, they require large datasets and raise concerns about explainability.

7.4 Fairness and Bias Mitigation Techniques

To address bias, techniques such as re-weighting training samples, removing sensitive attributes, and applying fairness constraints are employed. Post-processing methods adjust decision thresholds to reduce disparity between demographic groups.

8 System Architecture of AI-Driven Hiring

The AI-driven resume screening system follows a modular architecture designed to ensure scalability, transparency, and ethical compliance.

8.1 Architecture Overview

The system begins with resume ingestion, where resumes are uploaded through job portals or career websites. These resumes are parsed and converted into structured data. Feature extraction modules process both structured and unstructured data, generating inputs for the machine learning model.

The core scoring engine evaluates each candidate and assigns a relevance score based on job requirements. The shortlisted candidates are then reviewed by recruiters, ensuring human oversight. Audit logs and monitoring components track system performance and fairness metrics.



Figure 1: *High-level architecture of an AI-driven hiring and resume screening system*

The inclusion of a human-in-the-loop ensures that automated decisions are reviewed, maintaining accountability and trust.

9 Ethical Issues in AI-Driven Hiring

The use of Artificial Intelligence in hiring introduces several ethical challenges that must be carefully addressed to ensure responsible deployment. Hiring decisions have a direct impact on an individual's livelihood, career growth, and social mobility. Therefore, ethical lapses in automated hiring systems can cause significant harm to individuals and society.

One major ethical concern is the delegation of decision-making authority to algorithms. When hiring decisions are influenced or made by AI systems, accountability becomes

unclear. If a candidate is unfairly rejected, it may be difficult to determine whether responsibility lies with the recruiter, the organization, or the algorithm itself. This ambiguity raises concerns about moral responsibility and trust.

Transparency is another critical ethical issue. Many AI-driven hiring systems operate as black-box models, providing little explanation for why a candidate was shortlisted or rejected. Lack of transparency undermines trust and prevents candidates from understanding or contesting decisions that affect them.

Ethical hiring also requires respect for human dignity. Automated systems that reduce candidates to numerical scores may overlook contextual factors such as career breaks, non-traditional education paths, or personal challenges. Ethical AI systems must be designed to support human judgment rather than replace it entirely.

10 Privacy Concerns in Resume Screening

Privacy is a central concern in AI-driven resume screening because resumes contain sensitive personal data. This data may include names, addresses, contact details, educational history, employment records, and sometimes demographic information. Improper handling of such data can lead to serious privacy violations.

One major privacy issue is data collection without informed consent. Candidates may not be fully aware that their resumes are being processed by automated systems or used to train machine learning models. Ethical systems must clearly inform candidates about how their data will be used.

Data storage and retention also pose risks. Storing applicant data indefinitely increases the likelihood of data breaches and unauthorized access. Organizations must define clear retention policies and securely delete data once it is no longer required.

Another concern is secondary use of data. Resume data collected for hiring purposes should not be repurposed for unrelated analytics or shared with third parties without explicit consent. Ethical practice demands strict data minimization and purpose limitation.

Anonymization and encryption techniques are often recommended to protect candidate privacy. However, complete anonymization is challenging because certain resume attributes may indirectly reveal identity.

11 Bias and Discrimination in Automated Hiring

Bias in AI-driven hiring systems is one of the most widely discussed challenges in algorithmic decision-making. These systems often rely on historical hiring data, which may reflect existing social and organizational biases.

Gender bias is a common issue when historical data shows preference for male candidates in certain roles. Similarly, bias related to ethnicity, age, educational background, or geographic location may be embedded in training datasets. Even when explicit sensitive attributes are removed, models can learn proxy variables that correlate with protected characteristics.

Algorithmic bias can lead to discriminatory outcomes, where qualified candidates from underrepresented groups are systematically ranked lower. Such outcomes violate principles of fairness and equal opportunity and may also contravene employment laws.

Bias detection requires continuous auditing and subgroup analysis. Organizations must evaluate whether selection rates differ significantly across demographic groups and whether performance metrics are equitable.

Addressing bias is not a one-time task but an ongoing responsibility. Changes in job requirements, applicant pools, or data sources can reintroduce bias over time, necessitating regular monitoring.

12 Legal and Regulatory Considerations

AI-driven hiring systems must comply with employment laws and data protection regulations. Many jurisdictions prohibit discrimination based on protected characteristics such as gender, race, religion, age, or disability. Automated systems that produce biased outcomes may expose organizations to legal liability. Data protection regulations require lawful processing of personal data, transparency, and security. Candidates may have the right to know how their data is processed and to request deletion or correction of their information.

Another legal concern is automated decision-making. In some regulatory frameworks, individuals have the right to request human review of decisions made solely by automated systems. Therefore, organizations must ensure that AI tools support, rather than fully replace, human judgment.

Documentation and auditability are essential for legal compliance. Organizations should maintain records of model design, training data sources, evaluation metrics, and mitigation strategies to demonstrate due diligence.

13 Risk Analysis and Challenges

AI-driven hiring systems present multiple risks that must be proactively managed. These risks include technical, ethical, legal, and organizational challenges.

- **Bias Risk:** Models may unintentionally discriminate against certain groups.
- **Privacy Risk:** Data breaches or misuse of applicant data.
- **Explainability Risk:** Inability to justify hiring decisions.
- **Over-reliance on Automation:** Reduced human judgment and critical oversight.
- **Model Drift:** Degradation of model performance over time due to changing data patterns.

Managing these risks requires a combination of technical safeguards, governance policies, and human oversight.

14 Mitigation Strategies and Best Practices

To address ethical, privacy, and bias-related challenges, organizations should adopt a set of best practices.

Human-in-the-loop systems ensure that final hiring decisions involve human judgment, particularly in borderline cases. Fairness-aware model training and regular bias audits help detect and reduce discriminatory outcomes.

Privacy-by-design principles should be applied throughout the system lifecycle. This includes data minimization, encryption, access control, and clear consent mechanisms.

Transparency and explainability tools should be used to generate understandable reasons for decisions. This improves trust among candidates and recruiters.

Finally, organizations should establish governance frameworks involving legal, HR, and data science teams to oversee the responsible use of AI in hiring.

15 Results and Discussion

The implementation of an AI-driven hiring and resume screening system demonstrates significant improvements in efficiency and consistency when compared to traditional recruitment methods. Based on simulated evaluation and literature-supported out-

comes, the system was able to process a large volume of resumes in a short period of time, significantly reducing the workload of recruiters.

The automated screening model successfully identified candidates whose skills and experience closely matched job requirements. Recruiters reported improved shortlisting accuracy, as the AI system highlighted relevant candidates who may have been overlooked in manual screening processes. The use of semantic matching rather than simple keyword filtering enhanced candidate-job alignment.

However, the results also revealed challenges. Certain candidates with non-traditional career paths or employment gaps were ranked lower, indicating that the system may still favor conventional resume patterns. This reinforces the importance of human oversight and ethical safeguards.



Figure 2: *AI recruitment outcomes and Results*

Overall, the results indicate that AI-driven hiring systems can substantially support recruiters, provided that fairness checks, transparency mechanisms, and human-in-the-loop processes are implemented.

16 Advantages of AI-Driven Hiring

AI-driven hiring systems offer several advantages over traditional recruitment approaches:

- **Efficiency:** Automated screening significantly reduces the time required to review resumes.

- **Scalability:** AI systems can handle thousands of applications simultaneously.
- **Consistency:** Uniform evaluation criteria reduce random human bias.
- **Data-Driven Decisions:** Insights derived from analytics improve hiring strategies.
- **Cost Reduction:** Reduced recruiter workload lowers operational costs.

These advantages make AI-based recruitment particularly attractive for large organizations and high-volume hiring scenarios.

17 Limitations of AI-Driven Hiring

Despite its benefits, AI-driven hiring has notable limitations:

- **Algorithmic Bias:** Models trained on historical data may reinforce existing biases.
- **Lack of Explainability:** Complex models may not clearly justify decisions.
- **Privacy Risks:** Improper data handling can lead to privacy violations.
- **Over-Automation:** Excessive reliance on AI may reduce human judgment.
- **Data Dependency:** Poor-quality data results in inaccurate predictions.

These limitations highlight the need for responsible design, governance, and continuous evaluation.

18 Comparative Analysis of Hiring Approaches

Aspect	Traditional Hiring	AI-Driven Hiring
Resume Screening Time	High	Very Low
Scalability	Limited	Highly Scalable
Consistency	Varies by recruiter	High consistency
Bias Control	Subjective	Detectable and auditable
Transparency	Human reasoning	Requires explainability tools
Cost Efficiency	Moderate	High (long-term)

Table 1: Comparison between Traditional Hiring and AI-Driven Hiring

19 Conclusion

This case study has presented a comprehensive analysis of the application of Artificial Intelligence in hiring and resume screening, with particular emphasis on privacy, ethics, fairness, and accountability. As organizations increasingly adopt digital recruitment platforms and face growing volumes of job applications, AI-driven hiring systems have emerged as a practical solution to improve efficiency, scalability, and consistency in recruitment processes. The study demonstrated how machine learning and natural language processing techniques can automate resume screening, identify relevant candidate profiles, and support recruiters in making data-driven decisions. The findings of this case study highlight several advantages of AI-driven hiring systems. These systems significantly reduce time-to-hire, enable consistent application of evaluation criteria, and provide analytical insights that are difficult to achieve through manual screening. By leveraging semantic analysis and pattern recognition, AI models are capable of identifying suitable candidates beyond simple keyword matching, thereby enhancing the overall quality of shortlisting. When integrated effectively, such systems can improve recruiter productivity and organizational hiring outcomes.

However, this study also emphasizes that AI-driven hiring systems are not free from limitations and risks. Ethical challenges such as algorithmic bias, lack of transparency, and over-reliance on automation remain critical concerns. Models trained on historical data may unintentionally perpetuate existing social and organizational biases, leading to discriminatory outcomes. Furthermore, the opaque nature of complex machine learning models makes it difficult to explain or justify automated decisions, which can undermine trust among candidates and stakeholders. In conclusion, AI-driven hiring and resume screening systems hold significant potential to transform recruitment when designed and deployed responsibly. By integrating ethical principles, privacy safeguards, fairness-aware algorithms, and human oversight, organizations can harness the benefits of AI while minimizing risks. The responsible use of AI in hiring not only enhances operational efficiency but also promotes equitable and trustworthy recruitment practices aligned with societal values and ethical standards.

20 Future Scope

The future of AI-driven hiring and resume screening is highly promising and is expected to evolve continuously with advancements in artificial intelligence, data science, and ethical governance frameworks. As organizations increasingly rely on automated systems for recruitment, there is significant scope for improving the accuracy, fairness, and transparency of AI-based hiring solutions.

One important area of future development is the integration of explainable AI models that can provide clear and understandable justifications for hiring decisions. Such models will enhance transparency and help build trust among candidates, recruiters, and regulatory bodies. In addition, stronger bias detection and mitigation techniques can be developed to ensure fair treatment of candidates across different demographic groups.

Privacy-preserving machine learning approaches, such as federated learning and secure multi-party computation, offer promising solutions for protecting sensitive candidate data while still enabling effective model training. Future systems may also employ adaptive and continuously learning models that respond to changing job market trends, skill requirements, and organizational needs.

Furthermore, increased regulatory oversight and the establishment of standardized ethical frameworks are expected to shape the responsible use of AI in recruitment. Future research should focus on aligning technological innovation with ethical responsibility, legal compliance, and human-centered design to ensure that AI-driven hiring systems contribute positively to society and the workforce.



Figure 3: Future trends in AI-driven recruitment systems

21 References

1. Barocas, S., and Selbst, A. D., “Big Data’s Disparate Impact,” California Law Review.
2. Raghavan, M. et al., “Mitigating Bias in Algorithmic Hiring,” ACM Computing Surveys.
3. European Union, “General Data Protection Regulation (GDPR).”
4. IBM Research, “AI Fairness and Explainability in Hiring Systems.”
5. World Economic Forum, “Ethical AI in the Workplace.”